

First of all, we messed around with the different options for storing our metadata. We learned of the bitfield functionality in C, and decided to use that as we believe it's as small as we can get.

```
mbs189@ls:~$ ./a.out
size of int: 4
size of short: 2
size of unsigned int: 4
size of unsigned short: 2
sizeof "flags": 1
size of bitfield_test: 2
```

"bitfield\_test" wound up being the template for our metadata struct.

This was our design and strategy. Although, honestly it was a just a rough idea and the notes we took along the way. The following was written by Sakib as he implemented mymalloc() and myfree().

1. Initialize char myblock[4096]. Zero it. This block is the space we'll be managing.
2. Initialize a metadata struct inside this block that contains fields:
  - size - the amount of reservable memory
  - in\_use - whether or not the corresponding block of memory is reserved
  - prev - a pointer to the previous metadata struct or NULL
3. Say a user asks for x bytes:
  1. Check the in\_use field first block of memory. If it is in use, skip to the next block of memory (if it exists) and try again.
  2. If you find an unused block, compare its size to x:
    - If the block is too small, try the next block (if it exists).
    - If the block is too large, allocate the chunk you need and initialize the rest.
    - If the block is just right, allocate it.
4. Say a user asks to free a block of memory at address y.
  1. Start at address 0.
  2. Move forward the sizeof(metadata). If y is not equivalent to the current address, move forward the size of the metadata's corresponding block, and try again.
  3. If address y is found, then free its block, and concatenate it with any adjacent free blocks:
    1. Go backwards sizeof(metadata) indices.
    2. Set in\_use to 0.
    3. Go backwards prev\_size + sizeof(metadata) indices.
    4. If in\_use is 0, then set prev\_free to 1.
    5. Go forwards sizeof(metadata) + prev\_size + sizeof(metadata) + size.
    6. If in\_use is 0, then set next\_free to 1.
    - 7a. If prev\_free is 1 and next\_free is 0, then set the size of the previous block to prev\_size + size.

7b. If `prev_free` is 1 and `next_free` is 1, then set the size of the previous block to `prev_size + size + next_size`.

7c. If `prev_free` is 0 and `next_free` is 0, then do nothing.

7d. If `prev_free` is 0 and `next_free` is 1, then set the size of the current block to `size + next_size`.

Our results are pasted below for three test runs.

We were expecting Workloads C D and E to be much slower than A and B. That was not the case.

```
mbs189@ls:assign1$ ./memgrind
Mean Runtimes:
    workload A ---- 3.110000
    workload B ---- 24.030000
    workload C ---- 4.850000
    workload D ---- 3.520000
    workload E ---- 1.680000
    workload F ---- 0.040000
mbs189@ls:assign1$ ./memgrind
Mean Runtimes:
    workload A ---- 13.630000
    workload B ---- 107.800000
    workload C ---- 18.300000
    workload D ---- 15.320000
    workload E ---- 8.130000
    workload F ---- 0.150000
mbs189@ls:assign1$ ./memgrind
Mean Runtimes:
    workload A ---- 13.710000
    workload B ---- 110.020000
    workload C ---- 21.860000
    workload D ---- 20.810000
    workload E ---- 7.870000
    workload F ---- 0.090000
mbs189@ls:assign1$ _
```